

Kunal Aggrawal

+91-7417738058 | aggrawalk970@gmail.com | LinkedIn Profile | Github Profile

Education

University Of Delhi, Department of Computer Science	09/2023 – 08/2025
Master of Computer Applications (MCA), CGPA: 8.45	New Delhi, India
Sri Dev Suman University, Tehri Garhwal	07/2020 – 08/2023
Bachelor of Science (B.Sc.), CGPA: 7.87	Uttarakhand, India

Skills

Programming Languages: Python, SQL
Backend Development: FastAPI (REST API), Docker, PostgreSQL, SSMS, JWT, OAuth 2.0
AI/ML: Scikit-learn, Pandas, Numpy, Neural Networks, CNN, LangChain, LangGraph, Ollama
DevOps & Security: CI/CD, Nginx, VPS Deployment, SSH (key-only auth), Fail2Ban, EDR-aware server hardening
Tools & Platforms: VS Code, PyCharm, Jupyter Notebook, MySQL, Git, Postman

Experience

AI Engineer	Jan 2026 – Present
AmlgoLabs — Python, FastAPI, n8n, GitHub REST API, PostgreSQL, Nginx, LLM, OCR, CI/CD	
- Designed and deployed a GitHub repository automation system for a Bain & Company engagement using n8n and Python, identifying users inactive for 90+ days and revoking repository access via the GitHub REST API; engineered automated notifications and admin summary reports, reducing manual repo-hygiene effort to zero. - Built an end-to-end Car Comparison Dashboard integrating 7–10 OEM sources (Maruti, Kia, etc.) covering ~50 vehicle models, using AI-assisted pipelines to normalize and persist data to PostgreSQL for real-time queries; deployed on a VPS with Nginx and CI/CD for automated releases. - Developed an Invoice Data Extractor proof of concept for Suzuki, leveraging OCR and LLM-based parsing to extract structured fields (vendor, amounts, line items, dates) from a daily volume of ~3,000 invoices. - Delivered proof-of-concept solutions for 4 enterprise clients (Bain & Company, Suzuki, Lumax), including a supply-chain digitization POC for Lumax to replace manual Excel-based processes.	
Junior AI/ML Intern	Oct 2025 – Dec 2025
AmlgoLabs — Kubeflow, Kubernetes, Docker, MLflow, Python, Linux	
- Designed and built an on-premise ML lifecycle management proof of concept for Concentrix, covering model registry, serving, and monitoring using Kubeflow – first hands-on implementation of Kubernetes and Kubeflow. - Orchestrated containerized ML components via Kubernetes, ensuring secure and modular pipeline deployment within client infrastructure. - Implemented automated experiment-tracking workflows using MLflow for consistent logging and comparison of model runs. - Built 4 monitoring dashboards using Prometheus and Grafana for real-time model-serving visibility and alerting across deployed ML pipelines.	
Python Developer Intern	May 2025 – Sep 2025
Innovation HacksAI — FastAPI, SSMS, Autogen, REST API, Docker	
- Built and maintained scalable REST APIs using FastAPI for a Lead Generation Platform, improving response efficiency by 30%; implemented a RAG pipeline for personalized seller context, reducing manual analysis effort. - Integrated Google OAuth 2.0 and automated email sending, eliminating manual outreach entirely; added JWT-based authentication for secure route protection. - Integrated a Voice AI calling system (Vapi + Twilio) converting backend outreach content into live AI-driven phone calls, replacing manual email outreach with automated voice-agent calls for every product lead. - Engineered a Text-to-SQL multi-agent system using DeepSeek with a query-complexity orchestration layer and error-retry SQL generation, achieving ~80% query accuracy and producing automated data analysis, visualizations, and forecasting.	

Projects

AI Powered Blog Generator – Backend FastAPI, PostgreSQL, Google Gemini API	Github
- Developed a Blog Generation backend using FastAPI and Gemini AI API, automating content creation from user prompts and reducing manual effort by 90%; designed clean REST schemas with Pydantic for type-safe, consistent request handling. - Integrated PostgreSQL for persistent storage of user data and generated blogs, with modular DB handling built using Python classes for scalability.	
RAG-based Question Answering System Streamlit, Python, LangChain, Ollama	Github
Built a RAG-based QA system using LangChain and Ollama (Llama 3.2 1B), enabling context-aware querying across multiple scraped URLs via document chunking, vector embeddings, and similarity search – delivered through an interactive Streamlit UI.	

Certifications

Decode Data Science with Machine Learning PWskills Platform	Sep 2024
Covered foundational ML algorithms including regression, classification, and ensemble learning.	